

Explainable Deep Learning Framework for Pneumonia Detection and Classification Using Transfer Learning and Grad-CAM

Interim 1 Report

Department of Computing | BSc (Hons) in Computer Science

G.C.M.Sewwandi | 29002

Chapter 01 - Introduction

1.1 Chapter Overview

This chapter introduces the research project and establishes the context in which it is carried out. It opens with the clinical background of pneumonia diagnosis through chest X-ray imaging and the practical difficulties that motivate an automated and explainable approach. The problem is then framed at both a general and a discipline-specific level so that the precise research gap can be stated. From this gap the chapter derives a focused research question, explains the motivation behind the study, and sets out the aim and the objectives that direct the remaining work. A high-level picture of the proposed solution is presented to convey the intended workflow, followed by the hardware and software resources required, the boundaries of the project through an in-scope and out-of-scope definition, and a short summary that connects the chapter to the literature review that follows.

1.2 Problem Background

Pneumonia is one of the leading causes of illness and death worldwide and is responsible for millions of deaths each year, with children under five and elderly populations carrying a disproportionate share of that burden [1]. The disease is caused by a range of pathogens, including bacteria, viruses, fungi, and atypical organisms, and each of these presents with distinct radiographic and clinical characteristics that call for different treatment pathways. Distinguishing bacterial pneumonia, viral pneumonia, and COVID-19 from a normal lung is therefore not an academic exercise; it directly shapes antibiotic prescribing, isolation decisions, and the urgency of follow-up care.

Chest X-ray imaging remains the most widely adopted first-line modality for pulmonary disease because it is accessible, inexpensive, quick to acquire, and available even in low-resource settings. Reliable interpretation, however, is far from straightforward. The visual signatures of COVID-19, bacterial pneumonia, viral pneumonia, and a healthy lung frequently overlap, and the findings can be subtle or non-specific in early infection. Radiologist availability is uneven across institutions and regions, and the high patient volumes typical of emergency and intensive-care settings can introduce diagnostic delays that worsen outcomes. Against this backdrop, automated and explainable tools offer meaningful support by reducing radiologist workload, improving the consistency of preliminary assessments, and providing a rapid second opinion best positioned to augment, rather than replace, the clinician [1].

1.3 Problem Statement

1.3.1 General Problem

At a high level, the core problem is that timely and accurate differentiation of pneumonia subtypes from normal lung conditions is difficult to guarantee under real clinical pressures. When interpretation is delayed or incorrect, the consequences cascade: treatment is started late, antibiotics may be prescribed where they are not warranted, infectious cases may not be isolated promptly, and vulnerable patients face an elevated risk of deterioration. These pressures are amplified wherever radiological expertise is scarce or where patient throughput is high, which describes a large proportion of healthcare settings rather than a narrow exception. The general problem this research addresses is therefore the absence of a dependable, transparent, and readily deployable tool that can assist clinicians with rapid preliminary chest X-ray interpretation.

1.3.2 Specific Problem

Within the field of deep-learning-based chest X-ray analysis, several discipline-specific shortcomings narrow the general problem into a concrete gap. First, the majority of published studies treat the task as a binary problem normal versus abnormal, or COVID-19 versus non-COVID-19 rather than the clinically meaningful four-way distinction among Normal, COVID-19, Bacterial Pneumonia, and Viral Pneumonia [2], [3], [4]. Second, the large benchmark datasets that have driven progress, such as ChestX-ray8 [5] and CheXpert [6], are designed for multi-label thoracic abnormality recognition rather than focused pneumonia subtype classification. Third, although strong architectures exist, there are few controlled comparisons of DenseNet121, EfficientNetV2B3, and EfficientNetV2S trained on a single dataset with an identical pipeline, which makes it hard to attribute performance differences to the architectures themselves. Fourth, many models behave as black boxes; without interpretability, clinicians cannot confirm that a prediction is grounded in genuine pathology rather than in spurious image artefacts, and explainability, where it is offered at all, tends to be a detached visualisation rather than a feature woven into the clinical workflow. Finally, complete and deployable decision-support systems — those that combine secure access, role-based dashboards, historical prediction management, and automated reporting are largely absent from the published literature. The specific research gap this project addresses is the lack of an end-to-end, explainable, four-class chest X-ray classification system that is engineered for real clinical deployment.

1.4 Research Question

The study is guided by a single principal research question:

How can an explainable deep learning framework be designed, implemented, and deployed to accurately classify chest X-ray images into Normal, COVID-19, Bacterial Pneumonia, and Viral Pneumonia while providing clinically meaningful interpretability and practical, secure decision-support functionality?

This principal question is supported by three sub-questions:

- Which of the candidate transfer-learning architectures: DenseNet121, EfficientNetV2B3, or EfficientNetV2S offers the most suitable balance of accuracy, interpretability, and deployment efficiency for four-class chest X-ray classification?
- How can Grad-CAM explainability be integrated as a workflow-embedded feature rather than an isolated visualisation, so that each prediction is accompanied by trustworthy visual evidence?

- How can the trained model be delivered as a secure, mobile-friendly platform that supports prediction history, automated multilingual reporting, and a safeguard against untrustworthy inputs?

1.5 Research Motivation

The motivation for this work arises from a persistent disconnect between what deep learning can achieve in a research setting and what is actually usable at the point of care. A model that reports a high accuracy figure in a notebook does little for a clinician who needs a result that is explainable, auditable, and embedded in a usable interface. Two observations shaped the project. The first is that clinical adoption depends on trust, and trust depends on transparency: a prediction is only useful when the clinician can see why the model reached it and can check that reasoning against the visible anatomy. The second is that the supporting machinery around a model secure login, prediction history, structured reporting, and access on ordinary mobile devices is what turns a classifier into a tool a clinic could realistically use. Building a system that takes interpretability and deployment as seriously as predictive accuracy was the central driver of this research.

1.6 Research Aim

The aim of this research is to design, develop, and evaluate an explainable, end-to-end clinical decision-support system that classifies chest X-ray images into four diagnostic categories using transfer-learned convolutional neural networks, integrates Grad-CAM visual explanations directly into the diagnostic workflow, and delivers this capability through a secure, mobile-friendly web platform with prediction history, automated multilingual report generation, and safeguards that withhold predictions on inputs the model was never trained to interpret.

1.7 Research Objectives

To achieve this aim, the following research-specific objectives are defined:

- 1.7.1 To identify the clinical and technical requirements of explainable four-class chest X-ray classification by reviewing the relevant literature, surveying the publicly available chest X-ray datasets, and establishing the gaps that an end-to-end decision-support system must close.
- 1.7.2 To analyse three modern convolutional architectures DenseNet121, EfficientNetV2B3, and EfficientNetV2S together with appropriate preprocessing, augmentation, and class-imbalance strategies, and to analyse the explainability and deployment requirements that a clinical tool must satisfy.
- 1.7.3 To design and develop a complete system comprising a four-class transfer-learning inference pipeline with integrated Grad-CAM, a Next.js web frontend, a PHP REST backend secured with JSON Web Tokens, a Python inference microservice, a relational database, an out-of-distribution safeguard, automated trilingual (English, Sinhala, and Tamil) PDF reporting, and a platform-assistance chatbot.
- 1.7.4 To evaluate the trained models using accuracy, precision, recall, F1-score, and confusion-matrix analysis, to assess the clinical relevance of the Grad-CAM explanations, and to verify that the deployed system delivers its intended decision-support functionality.

1.8 Rich Picture of the Proposed Solution

The proposed solution, named LungVision AI, is a web-based lung disease detection platform with mobile-friendly, progressive web application style delivery. Figure 1 illustrates the intended end-to-end workflow. A registered user uploads a chest X-ray through the web interface. The backend forwards the image to a dedicated inference microservice, which preprocesses the image to the resolution and normalisation expected by the selected model, runs the convolutional network to obtain class probabilities, and produces a Grad-CAM heatmap that highlights the regions most responsible for the prediction. Before the result is trusted, a post-hoc out-of-distribution check decides whether the upload actually resembles the chest X-rays the model was trained on; if it does not, the diagnosis is withheld and the user is warned. The class scores, predicted condition, confidence, Grad-CAM overlay, and out-of-distribution verdict are returned to the backend, persisted in the database, and presented on the dashboard with confidence visualisation. Every prediction is stored as part of the patient's history, can be exported as a structured PDF report in English, Sinhala, or Tamil, and is accompanied by an optional chatbot that helps users navigate the platform. Access to all of this is governed by secure authentication and role-based permissions.




Figure 1. End-to-end workflow of the proposed LungVision AI decision-support platform.

1.9 Resource Requirements

1.9.1 Hardware

Model training requires a workstation equipped with a CUDA-capable GPU together with sufficient system memory and storage to hold the chest X-ray dataset and the resulting model checkpoints. For deployment and day-to-day inference the demands are far lighter: a standard server or development machine capable of running the Apache web server, the MySQL database, and the Python inference microservice is sufficient, and predictions can be served on the CPU. End users require only an ordinary desktop, tablet, or mobile device with a modern web browser, since all heavy computation takes place on the server side.

1.9.2 Software

The implementation draws on an open-source software stack. Model development uses Python with TensorFlow and Keras, served through a FastAPI inference microservice. The application backend is built on PHP 8 running under XAMPP with a MySQL database, exposing a REST API secured by JSON Web Tokens. The frontend is a Next.js 14 application written in TypeScript and styled with Tailwind CSS, using NextAuth.js for session handling and Recharts for confidence visualisation. The platform-assistance chatbot is powered by a hosted large language model accessed through a server-side key. Supporting tools include the Pillow imaging library for preprocessing, version control with Git, and phpMyAdmin for database administration.

1.10 Project Scope

The boundaries of the project are summarised in Table 1. The work concentrates on building and evaluating the four-class classification capability and the surrounding decision-support platform; it deliberately excludes activities, such as prospective clinical trials, that lie beyond the remit of an undergraduate research project.

In Scope	Out of Scope
Four-class classification of chest X-rays into Normal, COVID-19, Bacterial Pneumonia, and Viral Pneumonia.	Diagnosis of pulmonary conditions outside these four categories, or of diseases from other imaging modalities (e.g. CT, MRI).
Comparative evaluation of DenseNet121, EfficientNetV2B3, and EfficientNetV2S under a unified training pipeline.	Training of entirely novel network architectures from scratch without transfer learning.
Grad-CAM explainability and a post-hoc out-of-distribution safeguard integrated into the workflow.	Pixel-level lesion segmentation or automated radiology report writing.
A deployable web platform with secure authentication, role-based access, prediction history, and trilingual PDF reporting.	Integration with hospital information systems, PACS, or electronic health records.
Offline evaluation on a curated public chest X-ray dataset.	Prospective clinical trials and regulatory certification for autonomous diagnosis.

Table 1. Project scope defining the in-scope and out-of-scope boundaries of the research

1.11 Chapter Summary

This chapter established the clinical and technical context for the research. It described the burden of pneumonia and the difficulty of reliable chest X-ray interpretation, framed the general and specific dimensions of the problem, and identified the gap that motivates an end-to-end, explainable, four-class decision-support system. From this gap it derived the research question, the motivation, the aim, and four research-specific objectives, and it outlined the proposed LungVision AI solution along with its resource requirements and scope. The next chapter reviews the literature that underpins these design choices and substantiates the identified research gap.

Chapter 02 - Literature Review

2.1 Chapter Overview

This chapter reviews the body of work on which the proposed system rests and positions the research against it. Rather than cataloguing studies in isolation, the review weighs each contribution against the requirements of a deployable, explainable, four-class chest X-ray classifier and records where the present project agrees with, departs from, or extends the existing literature. The chapter begins with a conceptual map that explains how the review is organised. It then surveys the clinical and computational domain at a high level, compares existing systems and frameworks, and undertakes a deeper technological analysis of the algorithms, designs, and workflows that are relevant to the project. A reflective section closes the chapter by drawing the threads together and justifying the research gap.

2.2 Conceptual Map of the Literature

The reviewed literature is organised into three concentric layers that move from breadth to depth. The outermost layer, the domain overview, situates the work within medical imaging and deep learning and accounts for roughly a tenth of the discussion. The middle layer examines existing systems, frameworks, and designs for pneumonia and COVID-19 detection, comparing them critically and accounting for around a third of the discussion. The innermost and largest layer is a technological analysis that decomposes the field into its algorithmic, design, and workflow dimensions and accounts for the remaining majority of the review. Threaded through all three layers are two recurring concerns predictive performance and clinical trustworthiness and it is the tension between them that the reflection ultimately resolves into a statement of the research gap.

2.3 Domain Overview

Deep learning has become the dominant paradigm in medical image analysis. Convolutional neural networks learn a hierarchy of visual representations directly from data, progressing from low-level edges and textures to high-level structural patterns without the handcrafted feature engineering that earlier methods relied upon [7]. The foundational demonstration that deep networks can achieve strong image-recognition performance, given suitable architectures and training strategies, motivated a wave of clinical applications aimed at supporting diagnosis while easing the burden on specialists [1], [8]. Comprehensive surveys confirm that convolutional networks have been applied successfully to classification, detection, segmentation, and prognosis across medical imaging, while cautioning that careful methodology is needed to manage dataset bias, class imbalance, and spurious correlations [9].

Progress in chest X-ray analysis specifically has been accelerated by large public benchmarks. The ChestX-ray8 dataset established a hospital-scale, multi-label thoracic abnormality benchmark of more than one hundred thousand images across fourteen disease labels [5], and CheXpert extended this paradigm with uncertainty labels and expert comparison protocols [6]. During the pandemic, dedicated COVID-19 chest X-ray collections enabled the rapid development of automated screening models [10]. These resources are invaluable, yet their orientation towards multi-label thoracic findings means they do not directly serve the focused, mutually exclusive four-class distinction Normal, COVID-19, Bacterial Pneumonia, and Viral Pneumonia that is clinically meaningful for triage and treatment selection. This observation frames the rest of the review: the building blocks exist, but they have rarely been assembled into the specific configuration this project requires.

A further characteristic of the domain, evident across these datasets, is the heterogeneity and imbalance of the available data. Chest X-rays are acquired on different equipment, under different protocols, and from different patient populations, so they vary considerably in resolution, contrast, and intensity, and the disease classes are seldom evenly represented because they follow natural prevalence rather than a balanced experimental design. Both properties have practical consequences for any classifier trained on such data: without careful normalisation a model can latch onto acquisition artefacts rather than pathology, and without explicit handling of imbalance it can drift toward the majority classes. The literature on deep learning in medical imaging repeatedly stresses that methodological care around these issues is as decisive as the choice of network [9], a lesson that directly shapes the preprocessing and class-weighting decisions described in the methodology.

2.4 Existing Systems, Frameworks, and Designs

The most influential early system in this space is CheXNet, which used a densely connected network to achieve radiologist-level pneumonia detection on chest X-rays and demonstrated convincingly that deep models could match expert performance on a focused task [11]. CheXNet set the template for much of the work that followed, but it framed pneumonia detection as a binary problem and did not separate the pneumonia subtypes that carry different clinical implications. In my assessment, this is the recurring limitation across the field: strong results are reported on coarse label sets that do not map cleanly onto the decisions a clinician actually has to make.

The COVID-19 pandemic produced a large volume of detection systems built on transfer learning. Apostolopoulos and Mpesiana applied pretrained convolutional networks to automatically detect COVID-19 from X-ray images and reported encouraging accuracy with relatively small datasets [2]. Ozturk et al. proposed an end-to-end deep model for automated COVID-19 detection that also offered a degree of localisation [3], and Narin et al. compared several pretrained convolutional networks for coronavirus detection from chest X-rays [4]. Collectively these studies confirmed that transfer learning is an effective response to the scarcity of annotated medical data, but they share three weaknesses that are directly relevant to the present work. They are predominantly binary or three-class formulations that set COVID-19 against a small number of alternatives; they were trained and evaluated on heterogeneous datasets with differing preprocessing and evaluation protocols, which makes cross-study comparison unreliable; and they generally stop at the model, offering little in the way of a deployable, secured, and auditable system around it.

It is worth dwelling on why the binary framing that dominates these systems is a substantive limitation rather than a mere labelling choice. A model that separates abnormal from normal, or COVID-19 from everything else, leaves untouched the very distinction that most influences management: bacterial pneumonia typically warrants antibiotics, viral pneumonia and COVID-19 generally do not, and COVID-19 in particular carries isolation implications the others do not share. Collapsing these conditions into a single abnormal category therefore discards clinically actionable information. A four-class formulation restores that information, but it also makes the task harder, because the radiographic appearances of the three disease classes overlap far more with one another than any of them does with a healthy lung. This added difficulty is one reason the four-class problem remains under-explored, and it raises the premium both on architectures capable of capturing subtle, multi-scale differences and on evaluation that looks beyond headline accuracy to per-class behaviour.

Beyond single-model systems, ensemble strategies have been explored to push accuracy higher. Huang and Liao combined a stacking ensemble with an attention-enhanced EfficientNetV2 to classify multiple chest diseases, including COVID-19, and reported competitive performance on a multi-class problem [12]. This line of work is instructive because it demonstrates the viability of the EfficientNetV2 family for multi-class chest disease classification and so supports the decision to include EfficientNetV2 variants in the present comparison. At the same time, ensembles increase computational cost and complexity, which works against the deployment-oriented goals of this project; a single well-chosen architecture that can be served efficiently on modest hardware is a better fit for a tool intended to run alongside an ordinary web stack.

Taken together, the existing systems establish feasibility but leave a clear opening. None of the reviewed designs simultaneously delivers a clinically meaningful four-class distinction, a controlled comparison of competitive architectures under one pipeline, explainability that is embedded in the workflow rather than presented as an afterthought, and the deployment scaffolding authentication, role-based dashboards, prediction history, and automated reporting that a usable clinical tool demands. The proposed system is designed precisely around closing that combined gap rather than improving any single metric in isolation.

2.5 Technological Analysis

This section analyses the technologies underpinning the project across three dimensions: the algorithms and network architectures, the system and software design, and the end-to-end workflow. The aim is not only to describe these technologies but to justify, against the alternatives, why each was selected for LungVision AI.

2.5.1 Algorithmic Analysis

Modern medical image classification rests on a sequence of landmark architectural advances in computer vision. The introduction of residual learning resolved the degradation problem that had previously made very deep networks difficult to train; skip connections allowed networks of more than a hundred layers to be optimised stably and delivered substantial gains over earlier designs [13]. Earlier milestones deep convolutional networks with non-linear activations and dropout regularisation, multi-scale inception modules, and very deep networks built from small convolutional filters together established the representational capacity and the training practices that modern medical classifiers inherit. Visualisation studies that revealed how intermediate convolutional layers respond selectively to image patterns provided early insight into the learned feature hierarchy and foreshadowed the explainability techniques used today [14].

DenseNet builds on this lineage with dense connectivity, in which each layer receives the feature maps of all preceding layers within a dense block. This promotes feature reuse, improves gradient flow, and reduces parameter redundancy relative to conventional architectures [15]. These properties are well matched to chest X-ray analysis, where discriminating between opacity distributions, peribronchial thickening, airspace consolidation, and ground-glass opacities depends on combining information across multiple spatial scales. The effectiveness of densely connected networks in exactly this setting was demonstrated by CheXNet, which reached radiologist-level pneumonia detection [11], and it is this combination of fine-grained pattern capture and parameter efficiency that makes DenseNet121 a strong candidate for the present task.

The EfficientNet family approaches scaling from a different angle. Rather than enlarging depth, width, or resolution in isolation, it applies a principled compound-scaling rule that grows all three dimensions together using a fixed set of coefficients [16]. EfficientNetV2 refined this family by introducing Fused-MBConv blocks in the earlier stages and progressive learning with adaptive regularisation, which improved training efficiency without sacrificing accuracy [17]. EfficientNetV2B3 offers a particularly attractive balance of accuracy and computational cost for deployment-oriented applications, while the larger EfficientNetV2S allows an assessment of whether additional capacity and a greater proportion of Fused-MBConv blocks translate into meaningful gains. Including both a cross-family contrast (DenseNet versus EfficientNetV2) and a within-family contrast (EfficientNetV2B3 versus EfficientNetV2S) is what makes the comparison informative rather than arbitrary.

Because annotated medical images are costly and time-consuming to obtain, transfer learning is central to the algorithmic strategy. Pan and Yang formally framed transfer learning as the reuse of knowledge from a source domain to improve learning in a related target domain [18], and Tajbakhsh et al. provided the empirical counterpart, showing that fine-tuning pretrained networks is consistently more effective than training from scratch across a range of medical imaging tasks [19]. Although ImageNet [20] contains natural rather than medical images, its early convolutional layers learn generic edges, textures, and gradients that transfer well, leaving only the higher-level representations to be specialised through fine-tuning. This evidence directly informs the two-stage transfer-learning protocol adopted in the methodology.

Explainability completes the algorithmic picture. Class activation mapping introduced the idea of localising the discriminative regions that drive a network's decision [14], and Grad-CAM generalised this by using the gradients flowing into the final convolutional layer to produce class-discriminative heatmaps without any architectural modification [21]. For a clinical tool this is essential rather than cosmetic: surveys of explainable artificial intelligence in medicine stress that interpretability is a precondition for clinical acceptance and responsible deployment [22], and broader treatments of the subject argue that transparency is indispensable wherever decisions carry high stakes [23]. The consistent message that artificial intelligence should augment the clinician with interpretable, workflow-integrated outputs rather than replace expert judgement [1], [8] is the principle around which the explainability design of this project is built.

The progression from class activation mapping to Grad-CAM also clarifies a design choice that matters in practice. The original class activation mapping method localised discriminative regions but required a particular network structure, with a global average pooling layer feeding directly into the classifier, which constrained the architectures it could be applied to [14]. Grad-CAM removed that constraint by deriving its importance weights from gradients rather than from a fixed layer arrangement, so it can be attached to a pretrained DenseNet121 or EfficientNetV2 backbone without redesigning the network [21]. For a project that deliberately compares heterogeneous architectures this generality is precisely what is needed, because the same explanation mechanism can be applied uniformly across all three models and the explainability comparison can be kept as fair as the accuracy comparison. It is this architectural neutrality, as much as the visual quality of the heatmaps, that makes Grad-CAM the appropriate choice here.

2.5.2 Design Analysis

The design analysis turns from algorithms to the architecture of the system that surrounds them. A recurring shortcoming of the reviewed literature is that the model is treated as the end product; here it is treated as one component within a layered, service-oriented design. The chosen architecture separates concerns into three cooperating tiers: a presentation tier, an application and data tier, and a dedicated inference tier.

The presentation tier is a Next.js application written in TypeScript and styled with Tailwind CSS. A component-based framework with server-side rendering was preferred over a traditional multi-page application because it supports a responsive, app-like experience across desktop, tablet, and mobile devices while keeping the codebase maintainable. Session handling is delegated to NextAuth.js, and confidence scores are visualised with a charting library, giving clinicians an immediate sense of how decisive a prediction is. The application tier is a PHP REST backend running under XAMPP and backed by a MySQL relational database; it is responsible for authentication with JSON Web Tokens, role-based authorisation, persistence of scans and activity logs, and orchestration of the report and chatbot endpoints. Keeping this tier deliberately conventional makes the system easy to host on the ubiquitous LAMP-style stack found in many institutions.

The most important design decision is the separation of inference into its own Python microservice. Deep-learning inference depends on TensorFlow and Keras, which sit far more naturally in the Python ecosystem than in PHP. By exposing the trained models behind a small FastAPI service, the design keeps the heavyweight machine-learning dependencies isolated from the web application, allows the model to be scaled or replaced independently, and supports a pluggable pipeline in which DenseNet121, EfficientNetV2B3, or EfficientNetV2S can be selected according to deployment needs. This decoupling is precisely the kind of deployment-aware design that the surveyed single-model studies omit. The relational schema reflects the same philosophy: users carry a role that governs access, every scan records its predicted class, per-class confidence scores, the model used, the Grad-CAM image, and an out-of-distribution verdict, and an activity log captures actions for accountability.

2.5.3 Workflow Analysis

The workflow analysis examines how these components cooperate to turn an uploaded image into a trustworthy, reportable result. The pipeline begins with secure authentication and a role-appropriate dashboard. When an image is uploaded, the backend forwards it to the inference service, which applies the model-specific preprocessing, runs the network to obtain the four class probabilities, and generates the Grad-CAM overlay in a single pass. A critical addition, absent from the reviewed systems, is an out-of-distribution gate placed between prediction and presentation. Because a softmax classifier is closed-world and will always assign one of its four labels even to an image that is not a chest X-ray at all a post-hoc check is used to decide whether the input genuinely resembles the training distribution. When it does not, the diagnosis is withheld and the user is told why, which is the kind of safety behaviour a clinical tool ought to exhibit.

Once a result is accepted, the workflow persists it so that it becomes part of the patient's prediction history, available for later review and trend visualisation. From any stored result the user can generate a structured PDF report containing the patient details, the predicted condition, the confidence scores, and the Grad-CAM visualisation, and can do so in English, Sinhala, or Tamil a feature that directly addresses accessibility for the local clinical context and that is unusual among published systems. A platform-

assistance chatbot runs alongside the workflow to help users understand how to upload images and interpret results, while remaining explicitly outside the role of a medical advisor. The result is a workflow in which classification, explanation, safety checking, history management, and reporting are not separate experiments but stages of one coherent clinical process.

The reporting stage deserves particular emphasis because it is where the system reaches beyond the screen and into the clinical record, and because accessibility is rarely treated as a first-class concern in the literature. Generating the report in English, Sinhala, and Tamil acknowledges that a tool intended for the local healthcare context must communicate with clinicians and patients in the languages they actually use, and that a result which cannot be read and acted upon has limited value however accurate it may be. Producing the report on the server, with the predicted condition, the confidence scores, and the Grad-CAM overlay laid out in a consistent structure, also makes each prediction auditable after the fact rather than ephemeral. Coupled with the prediction history and the activity log, this turns the platform from a one-shot classifier into a system that keeps a coherent, reviewable account of the decisions it has supported the kind of accountability that responsible deployment of clinical artificial intelligence requires [8], [23].

2.6 Reflection

Reading the literature against the requirements of a deployable clinical tool brings a consistent pattern into focus. The field has convincingly established that deep convolutional networks, initialised through transfer learning, can detect pneumonia and COVID-19 from chest X-rays with impressive accuracy [11], [2], [3], [4], and it has produced both the architectures [13], [15], [16], [17] and the explainability techniques [21], [14] needed to do so responsibly. What it has largely not done is bring these pieces together. Five gaps recur across the recent and otherwise strong literature: most studies remain binary rather than addressing the clinically meaningful four-class distinction; that four-class distinction among Normal, COVID-19, Bacterial Pneumonia, and Viral Pneumonia is under-explored within a single unified framework; few studies compare DenseNet121, EfficientNetV2B3, and EfficientNetV2S under one dataset and pipeline; explainability is typically an add-on visualisation rather than a workflow-integrated feature linked to patient records; and deployment concerns such as secure authentication, role-based dashboards, prediction history, and automated reporting are frequently absent.

These five gaps, drawn from literature that is for the most part recent, jointly define the research gap this project addresses. The contribution of LungVision AI is therefore not a single new algorithm but the deliberate integration of a controlled four-class architecture comparison, workflow-embedded Grad-CAM explainability, an out-of-distribution safeguard, and a secure, mobile-friendly deployment with multilingual reporting into one end-to-end decision-support system. The methodology chapter that follows sets out how this integrated system is designed, built, and evaluated.

Chapter 03 - Methodology

This chapter sets out the research methodology and the concrete technical approach used to build and evaluate LungVision AI. It first establishes the philosophical and strategic framing of the research the paradigm, approach, strategy, and fact-collection mechanisms before mapping the execution of the work onto a structured design-science workflow. The detailed engineering methodology, covering the dataset, preprocessing, model training, explainability, and system deployment, is presented within that

workflow, followed by the project-management approach, the timeline, the ethical considerations, and a closing summary.

3.1 Research Paradigm

The research is conducted within a pragmatic paradigm that combines a positivist stance toward measurement with a design-science orientation toward building an artefact. The positivist element is appropriate because the central claims of the work that one architecture classifies chest X-rays more accurately than another, for instance are objective and empirically testable through controlled experiments and quantitative metrics. The pragmatic element acknowledges that the ultimate purpose of the research is not knowledge for its own sake but a working tool that solves a real clinical problem, which shifts the emphasis toward what demonstrably works in practice. This framing justifies an approach in which a system is designed and constructed and its components are then measured against objective criteria.

3.2 Research Approach

The study follows a primarily deductive approach embedded within design-science research. The literature yields clear, testable expectations that transfer learning will outperform training from scratch under data constraints, and that competitive architectures will differ in measurable ways on a unified pipeline and the experiments are designed to confirm or refute these expectations rather than to induce theory from open-ended exploration. Surrounding this deductive core is the constructive logic of design science: an artefact is specified, built, demonstrated, and evaluated, and the lessons from evaluation feed back into its refinement.

3.3 Research Strategy

Two complementary strategies are used. The first is controlled experimentation: the three candidate architectures are trained and tested under an identical dataset split, preprocessing regime, augmentation policy, optimisation configuration, and evaluation protocol, so that any observed difference in performance can be attributed to the architectures themselves rather than to confounding experimental conditions. The second is prototyping through design-science research: the trained model is wrapped in a complete, functioning software system whose components the inference microservice, the REST backend, the relational database, the web frontend, the explainability and safety layers, and the reporting and chatbot features are progressively built and demonstrated. Together the two strategies cover both the scientific question of which model is best and the engineering question of how to deliver it usefully.

3.4 Fact Collection Mechanisms

Because the research concerns image classification rather than human subjects, the primary facts are drawn from secondary data and from controlled experiments rather than from surveys or interviews. Three mechanisms are used. First, a structured review of the academic literature establishes the state of the art, the available architectures and explainability techniques, and the research gap. Second, publicly available, de-identified chest X-ray datasets supply the labelled images used to train and evaluate the models. Third, the experiments themselves generate quantitative evidence accuracy, precision, recall, F1-score, confusion matrices, and the qualitative Grad-CAM overlays which constitutes the principal empirical record against which the system is judged.

3.5 Research Methodology Execution Workflow

The execution of the research is organised as a design-science workflow whose stages map directly onto the activities of problem identification, relevance justification, comparative analysis, objective definition, design and development, and evaluation and communication. Table 2 summarises how each stage is addressed in this project, after which the design and development stage which carries the detailed engineering methodology is expanded in full.

Workflow Stage	How it is Addressed in this Research
Problem Identification	Establishing the difficulty of reliable, timely chest X-ray interpretation and the absence of an explainable, deployable four-class decision-support tool (Sections 1.2–1.3).
Relevance Justification	Grounding the problem in the clinical burden of pneumonia and the need for interpretable, workflow-integrated AI support [1], [8], confirming the work addresses a genuine need.
Comparative Analysis and Gap Justification	Critically comparing existing systems and architectures in the literature review to derive the five-part research gap (Sections 2.4–2.6).
Definition and Finalisation of Objectives	Translating the gap into the four research objectives of identifying, analysing, designing/developing, and evaluating the system (Section 1.7).
Design, Development, and Data Management	Preparing the dataset, preprocessing and augmenting it, training the three architectures with transfer learning, integrating Grad-CAM and out-of-distribution detection, and building the full platform (Section 3.5.5).
Evaluation and Communication	Measuring the models with standard metrics, assessing explanation quality, validating system functionality, and communicating the outcomes through this report (Section 3.5.6).

Table 2. Mapping of the design-science execution workflow stages to how each is addressed in this research.

3.5.1 Problem Identification

The first stage fixed the problem precisely: classifying chest X-rays into four mutually exclusive categories: Normal, COVID-19, Bacterial Pneumonia, and Viral Pneumonia in a way that is accurate, interpretable, and deployable. Identifying the closed-world behaviour of softmax classifiers as a safety risk was part of this stage and led to the inclusion of an out-of-distribution safeguard in the design.

3.5.2 Relevance Justification

The relevance of the problem is anchored in the clinical reality that pneumonia subtypes demand different treatment, isolation, and follow-up decisions, so that a tool capable of rapid, transparent preliminary discrimination has tangible value [1]. The justification also rests on the principle that decision support must be interpretable and embedded in the clinical workflow to be adopted at all [8], which the design takes as a binding constraint rather than an optional feature.

3.5.3 Comparative Analysis and Gap Justification

The comparative analysis carried out in the literature review distinguished the present work from prior systems and justified the research gap. By weighing CheXNet [11], the COVID-19 transfer-learning studies [2], [3], [4], and the EfficientNetV2 ensemble approach [12] against the requirements of a deployable four-class tool, this stage confirmed that no existing system combines a unified architecture comparison, embedded explainability, an input-safety safeguard, and a secured, reportable deployment.

3.5.4 Definition and Finalisation of Objectives

The identified gap was translated into the four research objectives stated in Section 1.7. Each objective is deliberately research-specific rather than generic: to identify the requirements and datasets, to analyse the candidate architectures and deployment needs, to design and develop the integrated system, and to evaluate it against quantitative and qualitative criteria. Finalising these objectives fixed the boundaries of the development and evaluation work.

3.5.5 Design, Development, and Data Management

This stage contains the core engineering methodology and is organised into the dataset, the preprocessing pipeline, the data splitting and augmentation policy, the model architectures, the transfer-learning strategy, the hyperparameter configuration, the Grad-CAM explainability integration, the out-of-distribution safeguard, and the system deployment architecture.

Dataset Description

The dataset comprises 9,208 chest X-ray images drawn from publicly available medical imaging repositories [5], [10], [11] and spanning the four diagnostic categories. As Table 3 shows, the dataset is imbalanced: the Normal and Bacterial Pneumonia classes are substantially larger than the COVID-19 and Viral Pneumonia classes. This reflects natural prevalence but requires explicit mitigation, through class-weighted loss and stratified sampling, to stop the classifier from favouring the majority classes. The images are stored in JPEG format and vary in resolution, contrast, and pixel intensity because they originate from different equipment, patients, and acquisition conditions, which is itself a reason for systematic preprocessing and normalisation.

Class	Number of Images
COVID-19	1,281
Normal	3,270
Viral Pneumonia	1,656
Bacterial Pneumonia	3,001
Total	9,208

Table 3. Class distribution of the chest X-ray dataset across the four diagnostic categories.

Data Preprocessing

Because the images vary so widely at source, a uniform preprocessing pipeline is applied across all subsets. Each image is resized to the resolution required by its architecture 224×224 pixels for DenseNet121 and 300×300 pixels for the EfficientNetV2 variants so that the comparison remains fair under consistent data partitions and architecture-specific normalisation. Since chest X-rays are

inherently grayscale while the ImageNet-pretrained backbones expect three-channel input, the single luminance channel is replicated across all three channels, allowing the pretrained weights to be reused without modifying the architecture. Pixel values are normalised using the scaling appropriate to each model. A data quality audit is performed before training to remove duplicate, corrupted, and below-threshold images, which guards against data leakage and ensures that the evaluation metrics reflect clean, representative inputs.

Data Splitting and Augmentation

The dataset is partitioned into training, validation, and test subsets using stratified random sampling so that the class distribution is preserved in every subset and performance estimates are not biased by the imbalance. The training set is used for weight optimisation, the validation set guides hyperparameter selection and early stopping, and the held-out test set is reserved solely for final, unbiased evaluation. Augmentation is applied only to the training set and is kept deliberately conservative: small random rotations, horizontal and vertical translations, zoom, brightness adjustment, and mild contrast variation simulate realistic acquisition variability, while horizontal flipping is used cautiously in view of anatomical laterality and any augmentation that could distort diagnostically meaningful structures is avoided. Validation and test images undergo only deterministic resizing and normalisation, so that the metrics describe performance on unmodified, clinically realistic inputs.

Model Architectures

Three architectures representing distinct design philosophies are evaluated. DenseNet121 implements dense connectivity, in which each layer connects directly to every subsequent layer within a dense block, promoting multi-level feature reuse, mitigating vanishing gradients, and retaining both fine-grained and abstract representations [15]; its final classification layer is replaced with a four-neuron softmax output while the pretrained backbone is preserved. EfficientNetV2B3 is selected for its compound-scaling strategy and training efficiency [17]; its backbone is fitted with a custom head comprising global average pooling, dropout, a dense intermediate layer, and a four-class softmax output. EfficientNetV2S, the larger variant in the same family, is included to test whether its greater capacity and larger proportion of Fused-MBConv blocks yield meaningful accuracy gains over EfficientNetV2B3, and its classifier is likewise replaced with an identical four-class softmax head. In the deployed system these models form a pluggable pipeline, with DenseNet121 serving as the default production model.

Transfer Learning and Hyperparameter Configuration

Transfer learning proceeds in two sequential stages designed to balance adaptation against regularisation. In the first stage the pretrained backbone is frozen and only the custom classification head is trained, allowing the randomly initialised head to settle on task-specific decision boundaries without destabilising the high-quality pretrained features. In the second stage, selected deeper backbone layers are progressively unfrozen and fine-tuned at a much lower learning rate, so that high-level representations adapt to chest X-ray characteristics while the generic low-level features are largely preserved. This protocol is justified by established transfer-learning theory [18] and by empirical evidence from medical imaging [19], and it is applied identically across all three architectures. All models are trained with categorical cross-entropy loss and the Adam optimiser, using a learning rate of 1×10^{-4} for head training and 1×10^{-5} for fine-tuning, a batch size of 32, and dropout in the classification head. Early stopping on validation loss and a ReduceLROnPlateau scheduler control the training length, class

weights computed from the training distribution counter the imbalance, and the best checkpoint by validation accuracy is retained.

Grad-CAM Explainability

Grad-CAM is integrated so that every prediction is accompanied by a visual explanation. For a given image, the gradients of the predicted class score with respect to the final convolutional feature maps are globally average-pooled to obtain importance weights; these weights are combined linearly with the corresponding activation maps and passed through a ReLU operation to retain only positively contributing features [21]. The resulting heatmap is upsampled to the original resolution and overlaid on the chest X-ray to show the regions most influential to the prediction. A clinically meaningful overlay is expected to highlight relevant lung parenchyma rather than image borders, labels, or extra-thoracic structures. Grad-CAM is treated as supportive visual evidence rather than a replacement for expert interpretation, and the heatmap is generated in the same inference pass as the prediction so that explanation is intrinsic to the workflow.

Out-of-Distribution Safeguard

Because a four-way softmax is closed-world and will confidently assign one of its labels even to an image that is not a chest X-ray, a post-hoc out-of-distribution safeguard is added to the inference service without retraining the classifier. The safeguard examines inference-time signals derived from the model's own forward pass principally the Mahalanobis distance of an image in the penultimate feature space relative to class-conditional Gaussians with a shared covariance, complemented by energy and maximum-softmax signals to judge whether an input genuinely resembles the training distribution. The decision statistics and thresholds are calibrated once on the model's in-distribution data and targeted at retaining the large majority of genuine chest X-rays. When an input is flagged as out-of-distribution, the predicted diagnosis is withheld and the user is warned, which prevents the system from issuing confident but meaningless results on invalid inputs.

System Deployment Architecture

The trained model is delivered through a layered, service-oriented platform. The frontend uses Next.js, TypeScript, Tailwind CSS, and NextAuth.js and provides secure login, role-based dashboards, confidence visualisation, image management, user-activity tracking, and prediction history with responsive, mobile-friendly access across devices. The backend is a PHP REST API secured with JSON Web Tokens and backed by a MySQL database under XAMPP, responsible for authentication, authorisation, persistence, and orchestration. Inference runs in a separate Python FastAPI microservice that loads the trained Keras models and performs preprocessing, prediction, confidence scoring, Grad-CAM generation, and the out-of-distribution check, returning its results to the backend for storage. Each prediction can be exported as a structured PDF report containing the patient details, predicted class, confidence scores, Grad-CAM visualisation, and timestamp in English, Sinhala, or Tamil, and a server-side, hosted-language-model chatbot assists users with the platform while remaining outside the role of a medical advisor. The relational schema records users and their roles, scans with their per-class scores, model used, Grad-CAM image, and out-of-distribution verdict, and an activity log for accountability. Table 4 lists the principal functional components of the resulting clinical decision-support system.

Component	Function
Image Upload	Accepts chest X-ray images through the web interface.
CNN Prediction Module	Classifies images into Normal, COVID-19, Bacterial Pneumonia, or Viral Pneumonia.
Confidence Visualisation	Displays class probabilities and prediction confidence.
Grad-CAM Explanation	Highlights the image regions influencing the prediction.
Out-of-Distribution Safeguard	Withholds the diagnosis when an input does not resemble a valid chest X-ray.
Prediction History	Stores previous scan results for patient management.
Trilingual PDF Report Generator	Produces structured medical reports in English, Sinhala, and Tamil.
Platform-Assistance Chatbot	Helps users navigate the platform without offering medical advice.
Web Dashboard	Supports user, scan, and result management.
Mobile-Friendly Delivery	Enables responsive, app-like access across devices.
Secure Authentication	Uses JSON Web Token authentication and role-based access control.

Table 4. Principal functional components of the LungVision AI decision-support platform and their roles.

3.5.6 Evaluation and Communication

Model performance is evaluated with accuracy, precision, recall, F1-score, and confusion-matrix analysis. Overall accuracy provides a global measure but is insufficient under class imbalance, so per-class precision, recall, and F1-score are emphasised; class-wise recall is treated as especially important because a missed pneumonia or COVID-19 case can have serious clinical consequences. The confusion matrix exposes which categories are most often confused and gives qualitative insight into the errors, while the Grad-CAM overlays are inspected to judge whether the model attends to clinically relevant lung regions. Final model selection therefore weighs accuracy, stable class-wise behaviour, explanation quality, and deployment feasibility together rather than accuracy alone.

Under this protocol the three architectures are trained and tested on identical data splits, so that their scores are directly comparable and any difference can be attributed to the architecture and its fine-tuning rather than to inconsistent data handling. The comparison considers not only the headline accuracy but also the per-class precision, recall, and F1-score, because a model that is strong on average may still be weak on a minority class that matters clinically, such as COVID-19 or viral pneumonia. Where two architectures perform within a narrow margin of one another, the decision is informed by factors beyond raw accuracy: the stability of class-wise behaviour, the inference cost and memory footprint relevant to mobile-friendly deployment, and the clinical plausibility of the Grad-CAM explanations. A marginally less accurate model whose heatmaps consistently fall on relevant lung regions may be preferred over a marginally more accurate one with diffuse or unreliable attention, since trustworthy explanations are central to the aims of this work. The scores, explanation quality, and

platform behaviour obtained under this protocol are reported and discussed in the subsequent evaluation, so that the final model is selected on the combined evidence rather than on accuracy alone.

3.6 Project Management Methodology

The project is managed using an iterative, Agile approach in the spirit of Scrum, which suits a research-and-development effort whose requirements are refined as understanding grows. The work is divided into short, focused iterations, each delivering a demonstrable increment for example, a working data pipeline, a trained model, the inference microservice, or the reporting feature so that risk is surfaced early and progress is continually visible. A lightweight backlog captures the tasks implied by the four objectives, and each iteration ends with a review of the increment against its acceptance criteria. This incremental delivery model was preferred over a rigid waterfall plan because the comparative experiments and the system build inform one another, and findings from one frequently reshape the priorities of the other.

3.7 Project Timeline

The project is sequenced into the phases shown in Table 5, beginning with the literature review and requirements analysis and ending with final evaluation and report writing. The phases overlap where it is productive data preparation continues into early model training, and system development proceeds alongside model refinement but the overall progression moves from understanding the problem, through building and training, to integrating, evaluating, and documenting the system.

Phase	Principal Activities
Phase 1 - Research and Requirements	Literature review, gap analysis, dataset selection, and finalisation of objectives.
Phase 2 - Data Preparation	Data quality audit, preprocessing pipeline, stratified splitting, and augmentation setup.
Phase 3 - Model Development	Two-stage transfer-learning training and comparative evaluation of the three architectures.
Phase 4 - Explainability and Safety	Grad-CAM integration and calibration of the out-of-distribution safeguard.
Phase 5 - System Development	Frontend, REST backend, database, inference microservice, trilingual reporting, and chatbot.
Phase 6 - Evaluation and Documentation	End-to-end testing, results analysis, and preparation of the interim and final reports.

Table 5. Project timeline showing the development phases and their principal activities.

3.8 Ethical Considerations

Several ethical considerations shape the work. All chest X-ray data is drawn from publicly available, de-identified repositories, so no personally identifiable patient information is collected or processed for model development, and no human-subject experimentation is involved. Within the deployed platform, user accounts are protected by hashed credentials, token-based authentication, and role-based access

control, and an activity log supports accountability. The most important consideration is the framing of the system itself: every report carries an explicit disclaimer that the platform is a decision-support tool and does not replace professional medical diagnosis, the out-of-distribution safeguard prevents confident predictions on invalid inputs, and the chatbot is constrained so that it assists with using the platform rather than offering clinical advice. The system is, by design, positioned to support rather than supplant the clinician, with final authority resting firmly with the medical professional.

3.9 Chapter Summary

This chapter set out the methodology of the research, from its pragmatic, design-science framing through to the detailed engineering of the system. It explained the research paradigm, approach, strategy, and fact-collection mechanisms, mapped the work onto a structured execution workflow, and detailed the dataset, preprocessing, augmentation, model architectures, two-stage transfer learning, Grad-CAM explainability, out-of-distribution safeguard, and deployment architecture. It also described the evaluation and model-selection protocol, the Agile project-management approach, the project timeline, and the ethical considerations that govern the work. Together these establish a reproducible and responsibly-constructed basis for the continued development and final evaluation of LungVision AI.

References

- [1] E. J. Topol, “High-performance medicine: The convergence of human and artificial intelligence,” *Nature Medicine*, vol. 25, no. 1, pp. 44–56, 2019.
- [2] I. Apostolopoulos and T. A. Mpesiana, “Covid-19: Automatic detection from X-ray images utilizing transfer learning with convolutional neural networks,” *Physical and Engineering Sciences in Medicine*, vol. 43, no. 2, pp. 635–640, 2020.
- [3] T. Ozturk et al., “Automated detection of COVID-19 cases using deep neural networks with X-ray images,” *Computers in Biology and Medicine*, vol. 121, p. 103792, 2020.
- [4] A. Narin, C. Kaya, and Z. Pamuk, “Automatic detection of coronavirus disease (COVID-19) using X-ray images and deep convolutional neural networks,” *Pattern Analysis and Applications*, vol. 24, no. 3, pp. 1207–1220, 2021.
- [5] X. Wang, Y. Peng, L. Lu, Z. Lu, M. Bagheri, and R. M. Summers, “ChestX-ray8: Hospital-scale chest X-ray database and benchmarks,” in *Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR)*, 2017, pp. 2097–2106.
- [6] J. Irvin et al., “CheXpert: A large chest radiograph dataset with uncertainty labels and expert comparison,” in *Proc. AAAI Conf. Artif. Intell.*, 2019, vol. 33, pp. 590–597.
- [7] Y. LeCun, Y. Bengio, and G. Hinton, “Deep learning,” *Nature*, vol. 521, no. 7553, pp. 436–444, 2015.
- [8] E. H. Shortliffe and M. J. Sepúlveda, “Clinical decision support in the era of artificial intelligence,” *JAMA*, vol. 320, no. 21, pp. 2199–2200, 2018.
- [9] G. Litjens et al., “A survey on deep learning in medical image analysis,” *Medical Image Analysis*, vol. 42, pp. 60–88, 2017.
- [10] M. E. H. Chowdhury et al., “Can AI help in screening viral and COVID-19 pneumonia?” *IEEE Access*, vol. 8, pp. 132665–132676, 2020.

- [11] P. Rajpurkar et al., "CheXNet: Radiologist-level pneumonia detection on chest X-rays with deep learning," arXiv preprint arXiv:1711.05225, 2017.
- [12] M.-L. Huang and Y.-C. Liao, "Stacking ensemble and ECA-EfficientNetV2 convolutional neural networks on classification of multiple chest diseases including COVID-19," IEEE Access, vol. 11, pp. 39243–39268, 2023.
- [13] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR), 2016, pp. 770–778.
- [14] B. Zhou, A. Khosla, A. Lapedriza, A. Oliva, and A. Torralba, "Learning deep features for discriminative localization," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR), 2016, pp. 2921–2929.
- [15] G. Huang, Z. Liu, L. Van Der Maaten, and K. Q. Weinberger, "Densely connected convolutional networks," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR), 2017, pp. 4700–4708.
- [16] M. Tan and Q. V. Le, "EfficientNet: Rethinking model scaling for convolutional neural networks," in Proc. 36th Int. Conf. Mach. Learn. (ICML), 2019, pp. 6105–6114.
- [17] M. Tan and Q. V. Le, "EfficientNetV2: Smaller models and faster training," in Proc. 38th Int. Conf. Mach. Learn. (ICML), 2021, pp. 10096–10106.
- [18] S. J. Pan and Q. Yang, "A survey on transfer learning," IEEE Transactions on Knowledge and Data Engineering, vol. 22, no. 10, pp. 1345–1359, 2010.
- [19] N. Tajbakhsh et al., "Convolutional neural networks for medical image analysis: Full training or fine tuning?" IEEE Transactions on Medical Imaging, vol. 35, no. 5, pp. 1299–1312, 2016.
- [20] J. Deng, W. Dong, R. Socher, L.-J. Li, K. Li, and L. Fei-Fei, "ImageNet: A large-scale hierarchical image database," in Proc. IEEE Conf. Comput. Vis. Pattern Recognit. (CVPR), 2009, pp. 248–255.
- [21] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, "Grad-CAM: Visual explanations from deep networks via gradient-based localization," in Proc. IEEE Int. Conf. Comput. Vis. (ICCV), 2017, pp. 618–626.
- [22] E. Tjoa and C. Guan, "A survey on explainable artificial intelligence (XAI): Toward medical XAI," IEEE Transactions on Neural Networks and Learning Systems, vol. 32, no. 11, pp. 4793–4813, 2021.
- [23] A. B. Arrieta et al., "Explainable Artificial Intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI," Information Fusion, vol. 58, pp. 82–115, 2020.